



Report

- Visualization of data values for the first row:



- Histogram of p values



Based on the histogram of p-values for gene-environment interactions, we can interpret the results as follows:

1. Distribution shape: The histogram shows a relatively uniform distribution of p-values across the range from 0 to 1, with a slight increase in frequency towards higher p-values.


2. Interpretation:

- P-values close to 0 indicate strong evidence against the null hypothesis (i.e., there is likely a significant gene-environment interaction).

- P-values close to 1 indicate weak evidence against the null hypothesis (i.e., there is likely no significant gene-environment interaction).

3.Findings: There is a slight increase in frequency for higher p-values, suggesting that for many genes, there may not be strong evidence of gene-environment interactions. However, there are still a considerable number of genes with low p-values, indicating potential significant interactions for some genes.

- Gene expression value for a row with a low p-value



✓ Data preparation

✓ Data Description

Data is generated from white blood cells from 48 individuals. Key details:

- A single file with 48 columns of data, plus some auxiliary columns
- Auxiliary columns: Probe name, Gene Symbol, Entrez Gene Id (ignore the rest)
- A single gene (identified by a Gene Symbol or Entrez Gene Id) could have multiple probes
- Total of 41,094 probes

Data Columns:

1. 12 Male Non-smokers (106-117)
2. 12 Male Smokers (118-129)
3. 12 Female Non-Smokers (130-141)
4. 12 Female Smokers (142-153)

Important notes:

- Values are logs to the base 2 of the original value
- Some 0 values are present due to thresholding low values before taking the log

```
import pandas as pd
import numpy as np
```

```
# Load the data
data = pd.read_csv('Raw Data_GeneSpring.txt', sep='\t') # Load data using tab as separator
```

```
# Display the first few rows and basic information about the dataset
print(data.head())
print(data.info())
```

```
➡ ProbeName GSM1149106_M_Ns3.txt:gProcessedSignal(normalized) \
0 GE_BrightCorner 17.288560
1 DarkCorner 2.172766
2 A_24_P66027 11.954556
3 A_32_P77178 6.224496
4 A_23_P212522 9.328137

GSM1149107_M_Ns4.txt:gProcessedSignal(normalized) \
0 18.103434
1 0.075925
2 11.681619
3 6.963845
4 9.164498

GSM1149108_M_Ns5.txt:gProcessedSignal(normalized) \
0 18.280110
1 0.232402
2 11.831579
3 7.096990
4 8.934509

GSM1149109_M_Ns6.txt:gProcessedSignal(normalized) \
0 17.883734
1 0.052573
2 12.262896
3 7.319873
4 9.166321

GSM1149110_M_Ns7.txt:gProcessedSignal(normalized) \
0 17.879555
1 0.483300
2 11.233925
3 6.611215
4 9.175541

GSM1149111_M_Ns8.txt:gProcessedSignal(normalized) \
0 17.340961
1 0.240482
2 11.722222
3 6.760409
4 8.810351

GSM1149112_M_Ns9.txt:gProcessedSignal(normalized) \
0 16.226210
```

```

1                0.327676
2                11.582652
3                7.810502
4                10.665818

GSM1149113_M_Ns10.txt:gProcessedSignal(normalized) \
0                16.203045
1                0.041980
2                13.039122
3                8.377619
4                10.681737

GSM1149114_M_Ns11.txt:gProcessedSignal(normalized) ... \
a                15.010000

# Separate auxiliary columns and gene expression data
auxiliary_columns = ['ProbeName', 'GeneSymbol', 'EntrezGeneID', 'Go']
gene_expression_columns = [col for col in data.columns if col not in auxiliary_columns]

# Create separate dataframes for auxiliary information and gene expression data
auxiliary_data = data[auxiliary_columns]
gene_expression_data = data[gene_expression_columns]

# Convert gene expression data to numeric type, replacing any non-numeric values with NaN
gene_expression_data = gene_expression_data.apply(pd.to_numeric, errors='coerce')

# Display basic statistics of the gene expression data
print(gene_expression_data.describe())

```

```

→ GSM1149106_M_Ns3.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.128693
std                  4.114148
min                  0.000000
25%                  2.932579
50%                  6.374927
75%                  9.182739
max                  18.298971

GSM1149107_M_Ns4.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.350905
std                  4.154255
min                  0.000000
25%                  3.195622
50%                  6.636967
75%                  9.374315
max                  18.557070

GSM1149108_M_Ns5.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.293946
std                  4.206224
min                  0.000000
25%                  3.023633
50%                  6.596164
75%                  9.382986
max                  18.698050

GSM1149109_M_Ns6.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.531112
std                  4.031424
min                  0.000000
25%                  3.623112
50%                  6.660977
75%                  9.434568
max                  18.612185

GSM1149110_M_Ns7.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.335960
std                  4.083656
min                  0.000000
25%                  3.301992
50%                  6.543399
75%                  9.283957
max                  18.552126

GSM1149111_M_Ns8.txt:gProcessedSignal(normalized) \
count                41093.000000
mean                 6.227665
std                  4.124976
min                  0.000000
25%                  3.113181
50%                  6.465138
75%                  9.219945

```

```
# Check for missing values
print("Missing values in gene expression data:")
print(gene_expression_data.isnull().sum().sum())
```

```
➞ Missing values in gene expression data:
0
```

```
# Group data by gender and smoking status
male_non_smokers = gene_expression_data.iloc[:, 0:12]
male_smokers = gene_expression_data.iloc[:, 12:24]
female_non_smokers = gene_expression_data.iloc[:, 24:36]
female_smokers = gene_expression_data.iloc[:, 36:48]
```

```
# Display the shape of each group
print("Shape of each group:")
print(f"Male Non-smokers: {male_non_smokers.shape}")
print(f"Male Smokers: {male_smokers.shape}")
print(f"Female Non-smokers: {female_non_smokers.shape}")
print(f"Female Smokers: {female_smokers.shape}")
```

```
➞ Shape of each group:
Male Non-smokers: (41093, 12)
Male Smokers: (41093, 12)
Female Non-smokers: (41093, 12)
Female Smokers: (41093, 12)
```

```
import matplotlib.pyplot as plt # Import matplotlib for plotting
import nbformat # Import nbformat to check its version
```

```
# Get the first row of gene expression data
first_row = gene_expression_data.iloc[0]
```


```
# Create a line plot of the first row values
plt.figure(figsize=(12, 6)) # Set figure size
plt.plot(range(len(first_row)), first_row.values, marker='o') # Plot values with markers
```

```
# Customize the plot
plt.title('Gene Expression Values for First Row') # Add title
plt.xlabel('Sample Index') # Label x-axis
plt.ylabel('Expression Value') # Label y-axis
plt.grid(True, linestyle='--', alpha=0.7) # Add grid lines
```

```
# Rotate x-axis labels for better readability
plt.xticks(range(len(first_row)), first_row.index, rotation=90)
```

```
# Adjust layout to prevent cutting off labels
plt.tight_layout()
```

```
# Show the plot
plt.show()
```



```
def custom_two_way_anova(data, factor1, factor2):
    n = len(data) # Total number of observations
    grand_mean = np.mean(data) # Overall mean

    # Calculate SST (Total Sum of Squares)
    sst = np.sum((data - grand_mean)**2) # Calculate total sum of squares

    # Calculate SS for Factor 1
    ss1 = sum(len(group) * (np.mean(group) - grand_mean)**2 for group in [data[factor1 == level] for level in np.unique(factor1)]) # Si

    # Calculate SS for Factor 2
    ss2 = sum(len(group) * (np.mean(group) - grand_mean)**2 for group in [data[factor2 == level] for level in np.unique(factor2)]) # Si

    # Calculate SS for Interaction
    ss_interaction = sum(len(group) * (np.mean(group) - grand_mean)**2 for group in [data[(factor1 == f1) & (factor2 == f2)] for f1 in ,

    # Calculate SSE (Error Sum of Squares)
    sse = sst - ss1 - ss2 - ss_interaction # Error sum of squares

    # Degrees of freedom
    df1 = len(np.unique(factor1)) - 1 # Degrees of freedom for factor 1
    df2 = len(np.unique(factor2)) - 1 # Degrees of freedom for factor 2
    df_interaction = df1 * df2 # Degrees of freedom for interaction
    dfe = n - (len(np.unique(factor1)) * len(np.unique(factor2))) # Degrees of freedom for error

    # Mean Squares
    ms1 = ss1 / df1 # Mean square for factor 1
    ms2 = ss2 / df2 # Mean square for factor 2
    ms_interaction = ss_interaction / df_interaction # Mean square for interaction
    mse = sse / dfe # Mean square error

    # F-values
    f1 = ms1 / mse # F-value for factor 1
    f2 = ms2 / mse # F-value for factor 2
    f_interaction = ms_interaction / mse # F-value for interaction

    # P-values
    p1 = 1 - stats.f.cdf(f1, df1, dfe) # P-value for factor 1
    p2 = 1 - stats.f.cdf(f2, df2, dfe) # P-value for factor 2
    p_interaction = 1 - stats.f.cdf(f_interaction, df_interaction, dfe) # P-value for interaction

    return (f1, p1), (f2, p2), (f_interaction, p_interaction) # Return F-values and p-values for both factors and interaction

def perform_two_way_anova(row_index):
    row_data = data.iloc[row_index] # Get the row data for the specified index

    # Prepare data for ANOVA
    male_non_smoker = row_data[male_non_smokers.columns].values # Extract male non-smoker values
```

```

male_smoker = row_data[male_smokers.columns].values # Extract male smoker values
female_non_smoker = row_data[female_non_smokers.columns].values # Extract female non-smoker values
female_smoker = row_data[female_smokers.columns].values # Extract female smoker values

# Combine all data into a single array
all_data = np.concatenate([male_non_smoker, male_smoker, female_non_smoker, female_smoker]) # Combine all data

# Create factor arrays
gender = np.repeat(['Male', 'Female'], len(all_data) // 2) # Create gender factor array
smoking = np.tile(np.repeat(['Non-smoker', 'Smoker'], len(male_non_smoker)), 2) # Create smoking factor array

# Perform two-way ANOVA using custom function
_, _, (, p_interaction) = custom_two_way_anova(all_data, gender, smoking) # Perform custom two-way ANOVA

return p_interaction # Return p-value for interaction (different response to smoke in men vs women)

# Create an array to store p-values for each row
p_values = np.zeros(data.shape[0]) # Initialize array with zeros for each row

# Iterate through each row and calculate p-value
for i in range(data.shape[0]):
    p_values[i] = perform_two_way_anova(i) # Calculate and store p-value for each row
    print(f"Processing row {i+1}/{data.shape[0]}, p-value: {p_values[i]:.4f}") # Print progress and p-value for each row

# Print the shape of the p_values array to confirm
print(f"Shape of p_values array: {p_values.shape}") # Display the shape of the resulting array

Processing row 1/41093, p-value: 0.9761
Processing row 2/41093, p-value: 0.5642
Processing row 3/41093, p-value: 0.4624
Processing row 4/41093, p-value: 0.6588
Processing row 5/41093, p-value: 0.8536
Processing row 6/41093, p-value: 0.5111
Processing row 7/41093, p-value: 0.7951
Processing row 8/41093, p-value: 0.5469
Processing row 9/41093, p-value: 0.8972
Processing row 10/41093, p-value: 0.9469
Processing row 11/41093, p-value: 0.5371
Processing row 12/41093, p-value: 0.6049
Processing row 13/41093, p-value: 0.6592
Processing row 14/41093, p-value: 0.3179
Processing row 15/41093, p-value: 0.5989
Processing row 16/41093, p-value: 0.8493
Processing row 17/41093, p-value: 0.3753
Processing row 18/41093, p-value: 0.8546
Processing row 19/41093, p-value: 0.2126
Processing row 20/41093, p-value: 0.3606
Processing row 21/41093, p-value: 0.7641
Processing row 22/41093, p-value: 0.5702
Processing row 23/41093, p-value: 0.6062
Processing row 24/41093, p-value: 0.7621
Processing row 25/41093, p-value: 0.7680
Processing row 26/41093, p-value: 0.7760
Processing row 27/41093, p-value: 0.8274
Processing row 28/41093, p-value: 0.2674
Processing row 29/41093, p-value: 0.9726
Processing row 30/41093, p-value: 0.7795
Processing row 31/41093, p-value: 0.1684
Processing row 32/41093, p-value: 0.2675
Processing row 33/41093, p-value: 0.9584
Processing row 34/41093, p-value: 0.7100
Processing row 35/41093, p-value: 0.2198
Processing row 36/41093, p-value: 0.7664
Processing row 37/41093, p-value: 0.9002
Processing row 38/41093, p-value: 0.5480
Processing row 39/41093, p-value: 0.7606
Processing row 40/41093, p-value: 0.8798
Processing row 41/41093, p-value: 0.1550
Processing row 42/41093, p-value: 0.6512
Processing row 43/41093, p-value: 0.2734
Processing row 44/41093, p-value: 0.0233
Processing row 45/41093, p-value: 0.4668
Processing row 46/41093, p-value: 0.1403
Processing row 47/41093, p-value: 0.3644
Processing row 48/41093, p-value: 0.7771
Processing row 49/41093, p-value: 0.3745
Processing row 50/41093, p-value: 0.5188
Processing row 51/41093, p-value: 0.3238
Processing row 52/41093, p-value: 0.4046
Processing row 53/41093, p-value: 0.3527
Processing row 54/41093, p-value: 0.6807
Processing row 55/41093, p-value: 0.7222
Processing row 56/41093, p-value: 0.3963
Processing row 57/41093, p-value: 0.8501
Processing row 58/41093, p-value: 0.6599

```

Note that the row 1 in data corresponds to the entry at index 0 in the p_values array.

```
import matplotlib.pyplot as plt
from matplotlib.widgets import Slider

# Create the main figure and axis
fig, ax = plt.subplots(figsize=(12, 6)) # Create a figure with a larger size for better visibility
plt.subplots_adjust(bottom=0.25) # Adjust the bottom margin to make room for the slider

# Plot the initial histogram
n, bins, patches = ax.hist(p_values, bins=50, edgecolor='black') # Create histogram with 50 bins and black edges


# Set labels and title
ax.set_xlabel('p-value') # Label for x-axis
ax.set_ylabel('Frequency') # Label for y-axis
ax.set_title('Histogram of p-values for Gene-Environment Interactions') # Title for the plot

# Create the slider axis
slider_ax = plt.axes([0.2, 0.1, 0.6, 0.03]) # Position of the slider [left, bottom, width, height]
slider = Slider(slider_ax, 'Row Range', 1, data.shape[0], valinit=data.shape[0], valstep=1) # Create slider object

# Function to update the plot when the slider is moved
def update(val):
    num_rows = int(slider.val) # Get the current value of the slider
    ax.clear() # Clear the current plot
    ax.hist(p_values[:num_rows], bins=50, edgecolor='black') # Plot new histogram with updated data
    ax.set_xlabel('p-value') # Reset x-axis label
    ax.set_ylabel('Frequency') # Reset y-axis label
    ax.set_title(f'Histogram of p-values for Gene-Environment Interactions (Rows 1-{num_rows})') # Update title with row range
    fig.canvas.draw_idle() # Redraw the figure

# Connect the update function to the slider
slider.on_changed(update) # Call update function when slider value changes

plt.show() # Display the plot
```



Based on the histogram of p-values for gene-environment interactions, we can interpret the results as follows:

1. Distribution shape: The histogram shows a relatively uniform distribution of p-values across the range from 0 to 1, with a slight increase in frequency towards higher p-values.
2. Interpretation:
 - P-values close to 0 indicate strong evidence against the null hypothesis (i.e., there is likely a significant gene-environment interaction).
 - P-values close to 1 indicate weak evidence against the null hypothesis (i.e., there is likely no significant gene-environment interaction).

3. Findings: There is a slight increase in frequency for higher p-values, suggesting that for many genes, there may not be strong evidence of gene-environment interactions. However, there are still a considerable number of genes with low p-values, indicating potential significant

interactions for some genes.

```
# Find rows with p-values < 0.05
significant_rows = data[p_values < 0.05] # Filter data for rows with p-values less than 0.05

# Print the number of significant rows
print(f"Number of rows with p-values < 0.05: {len(significant_rows)}") # Display count of significant rows

# Display the first few significant rows
print("\nFirst few significant rows:")
print(significant_rows.head()) # Show the first 5 rows of significant data

# Calculate the percentage of significant rows
percentage_significant = (len(significant_rows) / len(data)) * 100 # Calculate percentage of significant rows
print(f"\nPercentage of rows with p-values < 0.05: {percentage_significant:.2f}%") # Display percentage with 2 decimal places

# Optional: Save significant rows to a CSV file
significant_rows.to_csv('significant_gene_environment_interactions.csv', index=False) # Save significant rows to CSV
print("\nSignificant rows have been saved to 'significant_gene_environment_interactions.csv'") # Confirm save operation
```

➡ Number of rows with p-values < 0.05: 811

First few significant rows:

	ProbeName	GSM1149106_M_Ns3.txt:gProcessedSignal(normalized)	\
43	A_23_P111020	6.274315	
62	A_23_P151294	4.025540	
88	A_24_P926770	7.474671	
123	A_23_P145024	11.561870	
127	A_24_P910381	6.558103	

	GSM1149107_M_Ns4.txt:gProcessedSignal(normalized)	\
43	6.972955	
62	6.789796	
88	7.623500	
123	11.799589	
127	6.041059	

	GSM1149108_M_Ns5.txt:gProcessedSignal(normalized)	\
43	7.077835	
62	4.723725	
88	7.720977	
123	11.227537	
127	4.144054	

	GSM1149109_M_Ns6.txt:gProcessedSignal(normalized)	\
43	7.144897	
62	6.423915	
88	8.082088	
123	10.797735	
127	5.286856	

	GSM1149110_M_Ns7.txt:gProcessedSignal(normalized)	\
43	6.687301	
62	4.090033	
88	8.051618	
123	11.376524	
127	5.351962	

	GSM1149111_M_Ns8.txt:gProcessedSignal(normalized)	\
43	6.699010	
62	7.397929	
88	7.837849	
123	11.506655	
127	4.442767	

	GSM1149112_M_Ns9.txt:gProcessedSignal(normalized)	\
43	7.144043	
62	6.831710	
88	7.911806	
123	12.245893	
127	6.438065	

	GSM1149113_M_Ns10.txt:gProcessedSignal(normalized)	\
43	6.988871	
62	5.415668	
88	8.179203	
123	11.970752	
127	7.659210	

```
# Find the index of the minimum non-NaN p-value
min_p_value_index = np.nanargmin(p_values) # Get the index of the lowest non-NaN p-value in the numpy array

# Get the minimum non-NaN p-value
min_p_value = np.nanmin(p_values) # Get the minimum non-NaN p-value
```

```
# Print the row number (index + 1) and its p-value
print(f"Row number with minimum p-value: {min_p_value_index + 1}") # Print row number (index + 1 because indexing starts at 0)
print(f"Minimum p-value: {min_p_value:.6e}") # Print minimum p-value in scientific notation
```

```
➦ Row number with minimum p-value: 30369
Minimum p-value: 1.105540e-04
```

```
# Get the row 30369 of gene expression data
row_30369 = gene_expression_data.iloc[30368] # Subtract 1 because indexing starts at 0
```

```
# Create a scatter plot of the row 30369 values
plt.figure(figsize=(12, 6)) # Set figure size
```

```
# Define groups and their corresponding colors
groups = ['Male Non-smokers', 'Male Smokers', 'Female Non-smokers', 'Female Smokers']
colors = ['blue', 'red', 'green', 'purple']
```

```
# Plot each group separately
for i, (group, color) in enumerate(zip(groups, colors)):
    start = i * 12
    end = start + 12
    plt.scatter(range(start, end), row_30369.values[start:end], color=color, label=group) # Plot values with colored markers
```

```
# Customize the plot
plt.title('Gene Expression Values for Row 30369') # Add title
plt.xlabel('Sample Index') # Label x-axis
plt.ylabel('Expression Value') # Label y-axis
plt.grid(True, linestyle='--', alpha=0.7) # Add grid lines
```

```
# Move legend outside the plot area
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left') # Move legend outside and to the right
```

```
# Rotate x-axis labels for better readability
plt.xticks(range(len(row_30369)), row_30369.index, rotation=90)
```

```
# Add gene name or identifier as text annotation
gene_name = gene_expression_data.index[30368] # Get gene name or identifier
plt.text(0.02, 0.98, f"Gene: {gene_name}", transform=plt.gca().transAxes, verticalalignment='top', fontsize=10)
```

```
# Add p-value as text annotation
p_value = p_values[30368] # Get p-value for this gene
plt.text(0.02, 0.93, f"p-value: {p_value:.6e}", transform=plt.gca().transAxes, verticalalignment='top', fontsize=10)
```

```
# Adjust layout to prevent cutting off labels and legend
plt.tight_layout()
plt.subplots_adjust(right=0.85) # Make room for legend on the right
```

```
# Show the plot
plt.show()
```

```
# Print additional information about the gene
print(f"Gene: {gene_name}")
print(f"p-value: {p_value:.6e}")
print(f"Mean expression value: {row_30369.mean():.4f}")
print(f"Standard deviation: {row_30369.std():.4f}")
print(f"Minimum expression value: {row_30369.min():.4f}")
print(f"Maximum expression value: {row_30369.max():.4f}")
```

